

10 MARCH 2021

How to Add a Tenant View, OLAP View, or OLAPTS Table for Extract

5.21.11 - CreditLens™ Analytics and Reporting Configuration Guide

Contact

Moody's Analytics Support

MA_TechSupport@moodys.com

+1-212-553-1653

www.moodysanalytics.com/support

Documentation

Information Web: <https://information.moodysanalytics.com/>

Customer Portal: <https://moodys-analytics.force.com/customers/sfc/#search?tn=Content>

CONFIDENTIAL

This document contains information proprietary to Moody's Analytics, Inc. ("Moody's") and its affiliates. This information may not be reproduced, disclosed, or used in whole or in part without the express written permission of Moody's.

TRADEMARKS

Moody's Analytics, Moody's, and all other names, logos, and icons identifying Moody's Analytics and/or its products and services are trademarks of Moody's Analytics, Inc. or its affiliates. Third-party trademarks referenced herein are the property of their respective owners.

Table of Contents

To read a tenant view or join that view from the tenant directly to an olap table, such as `v_histstmtbalancelatest`, you can add the table directly into olap without causing any ETL transformation to take place. These tables can be added and flagged as system tables and set to `isactive_=false`.

You can run this query to see what tables are inactive in the `olap_etl_control` table and are also flagged as system tabletypes.

Query:

codeblock

```
select * from olap.olap_etl_control where sourcetable_ ilike 'v\_%' and etlfunction_ = 'olap_
metadata' ;

-- as of 5.21.11 you can run the following view below, which gives the same information as the previous
query --
select * from v_olap.latestviews;
```

Results:

	tablename_ [PK] character varying	tabletype_ character varying	isactive_ boolean	sequence_ integer	sourcetable_ character varying	etlfunction_ character varying	etlfuncio text	createddate_ timestamp without time zone
1	facthiststmtbalancelatest	fact	false	424	v_histstmtbalancelatest	olap_metadata		2020-06-25 03:28:41.059764
2	dimaddedaccountlatest	dimension	false	425	v_addedaccountlatest	olap_metadata		2020-06-25 03:28:41.059764
3	dimhiststmtlatest	dimension	false	431	v_histstmtlatest	olap_metadata		2020-06-26 17:54:41.431815
4	dimhiststmtconstantlatest	dimension	false	432	v_histstmtconstantlatest	olap_metadata		2020-06-26 17:54:41.431815

Complete the following procedures to create a new view inside of `tenantdb` called `v_histstmtlatest`, and expose it to the olap database for querying and data extract.

Query: Create a view in `tenantdb`.

codeblock

```
-- tenant prefix is not needed but prevents you from creating the view in olap accidentally --

DROP VIEW IF EXISTS tenant.v_histstmtlatest;
analyze tenant.historicalstatement;

create or replace view tenant.v_histstmtlatest as
select
  h.pkid as v_histstmtlatestid -- create pkid and rename it as viewname||id
  -- alternatively you can use h.id as primarykey and rename it same as viewname||id ---
  -- alternatively you can use h.id ||h.versionid as primarykey or use any set of columns to set
  granularity of table such as pkid ||accountid and name as viewname||id ---
  ,h.pkid :: varchar as pkid -- pull pkid for transformation if this is latestversion view
  ,h.id :: varchar as historicalstatementid_ -- pull the id_ from tenant table and name it as
  tenant.tablename||id_

-- add several jsonb columns here and alphabetize them in lowercase order
  ,(h.jsondoc ->> 'AccountId')::numeric as accountid
  ,(h.jsondoc ->> 'ColumnA')::text as columna
  ,(h.jsondoc_ ->> 'ColumnZ')::text as columnz

-- add all base table context attributes you want for your view here
```

```

,h.wfid :: varchar
,h.taskid :: varchar
,h.versionid ::int4
,h.isdeleted ::boolean
,h.islatestversion ::boolean
,h.baseversionid ::int4
,h.contextuserid :: varchar
,h.isvisible ::boolean
,h.isvalid ::boolean
,h.snapshotid ::int4
,h.t :: varchar
,h.createdby :: varchar
,h.createddate :: timestamp
,h.updatedby :: varchar
,h.updateddate_:: timestamp

-- add these 2 columns to be used for incremental transformation and incremental dataextracts --
, ( case when h.updateddate_> h.createddate_ then h.updatedby_ else h.createdby_ end ):: varc
har as sourcepopulatedby
, GREATEST( h.createddate_,h.updateddate_):: Timestamp as sourcepopulateddate_

-- add your FROM clause here - if using isvisible_, isvalid_, islatestversion_ and not isdeleted_ you can
also add the word ONLY to the FROM clause --
from only historicalstatement h
where ( h.isvalid and h.isvisible and h.islatestversion_ )
; -- note: and not h.isdeleted_ this was removed in 20.31.1

```

Important Notes on View Design:

- The view must be named as v_someview. This tells the olap code that it is a view that could potentially be converted as a table using the same format.
 - For specific customer views, name it as v_{customerinitials}viewname such as v_abctenantfinancialview. This prevents any views Moody's creates from overriding existing customer custom views
- The view must have the first column as v_someview+id_, such as v_someviewid_, which holds the granularity of the table. It is the primary key and cannot have duplicates.
- The view must have pkid_ separately, as pkid_ translates back to the latestversion of data.
- The view must have a column called id_ that is renamed as sourcetablename+id_. In this case, it is the historicalstatementid_. This allows joining back to the tenant source table using this id_ value.

Results:

Data Output
Messages
Explain
Notifications
Scratch Pad

CREATE VIEW

Query returned successfully in 249 msec.



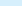



Query: Select from the view to ensure that it gives the data you want.

codeblock

```
-- tenant prefix is not needed, but prevents you from creating the view in olap accidentally --
```

```
select * from tenant.v_histstmlatest limit 5
```

Results:

Data Output											Messages	Explain	Notifications	Scratch Pad								
	v_histstmlatestid_ character varying		pkid_ character varying		historicalstatementid_ character varying		accountid_ numeric		columna_ text		columnz_ text		wfid_ character varying		taskid_ character varying		versionid_ integer		isdeleted_ boolean		islatestversion_ boolean	
1	83		83		83/2		[null]		[null]		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...				2	false		true
2	90		90		90/5		[null]		[null]		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...				5	false		true
3	89		89		89/10		[null]		[null]		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...				10	false		true
4	1		1		1/3		[null]		[null]		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...				3	false		true
5	2		2		2/3		[null]		[null]		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...				3	false		true

Switch user/database and log in as olap user in olap database

Query: This query create a view reference in the FDW for olap to query and also create a new table called olapts.facthiststmlatest.

codeblock

```
select olapts.populate_olapmetadata_views( 'v_histstmlatest' , 'dim' );
```

Results: The view is now registered as the same view in the olap database and is also registered as an olapts table if needed. This table will be set to inactive by default so as not to slow down the current Time Series job.

Query:

codeblock

```
select * from olapts.olap_etl_control where sourcetable_ ilike 'v\_%' and etlfunction_ = 'olap_  
metadata' ;
```

```
-- as of 5.21.11, you can run this view below, which gives same information as the previous query --  
select * from v_olapts_latestviews;
```

Results: Views are available to query in olap but will not generate an olapts table unless they are set to **inactive_=true**.

Note If **inactive_=true**, the view will materialize into an olapts table using existing DDL at that time. If the tenant.view is updated later, then this process needs to repeat to generate an updated table format in olapts.

Data Output								Messages	Explain	Notifications
	tablename_ [PK] character varying	tabletype_ character varying	inactive_ boolean	sequence_ integer	sourcetable_ character varying	etlfunction_ character varying	etlfunctionparameters_ text	createddate_ timestamp without time zone		
1	facthiststmbalancelatest	fact	false	907	v_histstmbalancelatest	olap_metadata		2020-08-07 05:01:07.864776		
2	dimaddedaccountlatest	dimension	false	908	v_addedaccountlatest	olap_metadata		2020-08-07 05:01:07.864776		
3	dimhiststmlatest	dimension	false	909	v_histstmlatest	olap_metadata		2020-08-07 05:01:07.864776		
4	dimhiststmlconstantlatest	dimension	false	910	v_histstmlconstantlate...	olap_metadata		2020-08-07 05:01:07.864776		
5	dimpeeranalysislatest	dimension	false	952	v_peeranalysislatest	olap_metadata		2020-08-17 04:47:57.774441		
6	factratingscenarioblockdata...	fact	false	953	v_ratingscenarioblockd...	olap_metadata		2020-08-17 04:47:57.774441		

Following are reasons to materialize a view into an olapts table:

- The view has too many rows or takes too long to run in real-time.
- You are unable to generate the data extract file directly from the view.
- The trade off is that the time that it takes executing queries against the view prior to being materialized will be the same time it takes to materialize a view during the next Time Series Job run.
- You are creating an olapts view called "view1" and it joins this view above (for example, v_histstmtlatest to olapts.dimfinancial, olapts.dimentity. - joining the results of this view to another olapts tables.
 - Joining an olapts view to olapts tables causes the newly created view ("view1" in previous bullet) to drop from existence or disappear from the olapts database on each run of the TimeSeries ETL process.

Note the following differences in data extracts:

- Extracts of a view are in real-time, as they come directly from the olap link into the tenantdb called Foreign Data Wrapper (FDW).
- Extracts of a materialized olapts table are a snapshot of data from the last run of the Time Series Job.

To materialize a view into an olapts table complete the following query:

Query: Update the row in olapts.olap_etl_control where sourcetable_='v_yournewviewname'.

codeblock








```
update olapts.olap etl control set isactive_ = true where sourcetable_ = 'v_histstmtlatest' and
etlfunction_ = 'olap_metadata' ;

-- as of 5.21.11, you can run the following view, which gives the same information as the previous query --
update v_olapts_latestviews set isactive_ = true where sourcetable_ = 'v_histstmtlatest' ;

-- or to enable all views at once, such as all IPRE views, you can use something like this statement

update olapts.olap etl control set isactive_ = true where sourcetable_ ilike 'v\_%pre%' and
etlfunction_ = 'olap_metadata' ;
-- as of 5.21.11, you can run the following view, which gives same information as the previous query --
update v_olapts_latestviews set isactive_ = true ;
```

Results: This view materializes into an olapts table on the next Time Series job run.

Data Output		Messages	Explain	Notifications	Scratch Pad		
	 tablename_ [PK] character varying	 tabletype_ character varying	 isactive_ boolean	 sequence_ integer	 sourcetable_ character varying	 etlfunction_ character varying	 etlfunctionparameters_ text
1	dimhiststmtlatest	dimension	true	422	v_histstmtlatest	olap_metadata	
2	facthistmtbalancelatest	fact	true	416	v_histmtbalancelatest	olap_metadata	
3	dimaddedaccountlatest	dimension	true	420	v_addedaccountlatest	olap_metadata	

The best practice is to enable one view at a time, then execute the TimeSeries Job, then check for errors. If the first view is correct, then enable the next one until all are enabled.

Execute the Reporting Database Time Series Job, and check for errors.

If you find any errors, make note of what the error message is, then go in and set the specific view/table back to **inactive** so that the job does not keep erroring out until you can fix your view.

codeblock

```
update olapts.olap etl control set isactive_ = false where sourcetable_ = 'v_histstmlatest' and
etlfunction = 'olap metadata' ;
-- as of 5.21.11, you can run the following view, which gives the same information as the previous query --
update v_olapts_latestviews set isactive_ = false ;
```

If you are creating a brand new tenant view to use in this process, you will need to analyze the view design and set it up so that there are no duplicates in the PrimaryKey, and the granularity of the view is such that it allows for only one unique row for each PrimaryKey regardless of version.

The Primary Key will normally be the id_ of the underlying table; however, if you have multiple tables involved, you must add key values from each table into a combined PrimaryKey for this view.

For example, if the view joins Entity, HistoricalStatement, Financial tables, then the PrimaryKey would be like this: (entity.pkid_||':'||historicalstatement.pkid_||':'||financial.pkid_) as v_viewnamepkid.

When using pkid_ as the primary key, you also need to have the WHERE entity.islatestversion_ and historicalstatement.islatestversion_ and financial.islatestversion_ filters - to remove duplicate versions from a materialized table that joins entity,historicalstatement,financial tables together.

If you receive a Success message, then check the new table by selecting from it. In this example, only 20 rows were selected.

Query:

codeblock

```
select * from olapts.dimhiststmlatest limit 20; -- this is your view called v_histstmlatest converted
into an olapts table.
```

Results: table output

Data Output		Messages	Explain	Notifications	Scratch Pad								
	dimhistmtlatestid_ [PK] character varying		pkid_ character varying		historicalstatementid_ character varying		accountid numeric		wfid_ character varying		taskid_ character varying		versionid_ integer
1	1		1		1 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
2	2		2		2 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
3	6		6		6 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
4	7		7		7 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
5	43		43		43 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
6	4		4		4 2		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		2
7	5		5		5 2		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		2
8	13		13		13 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
9	8		8		8 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
10	9		9		9 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
11	10		10		10 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
12	11		11		11 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1
13	136		136		136 1		[null]		00000000-0000-0000-0...		00000000-0000-0000-0...		1

Now, each time you run the Time Series job, the table will run an UPSERT and either Insert or Update any incremental rows from the view that have changed in tenant, and push them into your olapts materialized table.

This process allows for faster subsequent runs of the TimeSeries job using incremental updates to this new table.

Want to find out more? You can find the most up-to-date product documentation on Information Web (<https://information.moodyanalytics.com>).